We used Gemini to generate the images. First we generated an image of a man doing the dishes, with the prompt “could you generate a picture of a man doing the dishes”, which resulted in a really good picture of just that. When inspecting the picture, there are no obvious inconsistencies.

The second picture was generated with the prompt “could you generate a picture of a happy family”, which resulted in a picture of a family celebrating something outside. Again, it is very hard to find inconsistencies, the only things that look weird is the text on the cake, where it seems to say “14APPY PAMILY” and an odd looking object in the corner of the table cloth. The grandmother's hands could also be an inconsistency, but it is hard to tell. Why these things are inconsistent is probably because this is a more detailed picture, with more things in focus, resulting in a higher risk of mistakes when generating the picture. The text on the cake is still made out of letters, but they do not really make sense, which may be due to the fact Gemini does not take words and letters into consideration when generating the picture.

The third picture was generated with the prompt “could you generate a picture of two persons doing a cartwheel”. This picture also lacked inconsistencies, even though one can always argue about the physics of things, for example why the sand is in the way it is, why the clothes look like they do etc.

The fourth picture was generated with the prompt “could you generate a person holding 2 children”. When inspecting this picture, there are no direct inconsistencies. The logos on some of the clothes are a little bit unclear but other than that there is nothing unusual. The only other remarkable thing is that it says “ORMBERGET NATURRESERVAT” on a sign, and that Ormberget is a place in Luleå. This means that Gemini does not only take the prompt into consideration when generating the picture but also things like one’s location.

The next picture was generated with the prompt “could you generate a picture of three persons walking 2 dogs and one cat”. This picture was also generated really good, with few things that stood out. What can be mentioned is that the dogs seem to only have three legs, but it is something that can occur when taking a real picture as well, and that the handles for the leashes look quite strange. Other than that everything looks normal.

The final picture was generated with the prompt “ could you generate a picture of a zoo with a lot of animals and people”. This is the first picture that really has direct inconsistencies. At first glance the picture looks really good, but when zooming in it is clear that it is not a real picture. There are both inconsistencies in the picture but also in logic. For example, almost all of the signs are spelled wrong, except for the big “Välkommen till Luleå Zoo”-sign. When looking to the bottom left corner, it looks like there are kids inside of the enclosure, but the kids do not have any distinct features, and look more like some color blobs. The animals also look quite strange, and in the enclosure in the middle it looks like there are bears and elks in the same enclosure. Why these inconsistencies appear is probably because the picture has so many details, and because Gemini does not take logic or facts into considiration.

The conclusion is that gemini is really good at generating pictures in general, but when asked to generate detailed pictures some inconsistencies can appear. Text is also something it struggles with, and it is a little alarming that it takes one's location into consideration when generating the pictures. Something that almost all pictures have in common is that the lighting and quality almost always looks the same, and the people in the pictures rarely have any imperfections. It almost looks like a social media filter over them, to signal happiness and that everything is perfect.